

# Deep Reinforcement Learning with Double Q-learning

Friday, February 03, 2017 12:05

## Background

- The popular Q-learning algorithm is known to overestimate action values under certain conditions
  - because it includes a maximization step over estimated action values, which tends to prefer overestimated to underestimated values
  - Its irrespective of the source of approximation error (i.e. not because of the neural net or other approximation method)
- This paper answers the following:
  - whether, in practice, such overestimations are common,
  - whether they harm performance,
  - whether they can generally be prevented
- They adapt the Double Q-Learning algorithm to DQN and propose a specific adaptation
- Regular DQN:
  - We want to be able to estimate  $Q(s,a)$  - the true value of action  $a$  in state  $s$ 
$$Q_{\pi}(s, a) \equiv \mathbb{E}[R_1 + \gamma R_2 + \dots \mid S_0 = s, A_0 = a, \pi]$$
  - Given optimal  $Q$  we can easily find the policy
  - The standard Q learning update for the parameters after taking action  $A_t$  in state  $S_t$  and observing the reward  $R_{t+1}$  and resulting state  $S_{t+1}$  is given by:

$$\theta_{t+1} = \theta_t + \alpha(Y_t^Q - Q(S_t, A_t; \theta_t)) \nabla_{\theta_t} Q(S_t, A_t; \theta_t) \quad (1)$$

Where the target is defined as:

$$Y_t^Q \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t)$$

- In DQN paper they introduced 2 interesting ideas:
  - Fixed target - holding another network whose parameters were updated from the online network in a lower frequency. This network was used to calculate the target :
$$Y_t^{\text{DQN}} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t^-)$$
  - Experience Replay - observed transitions are stored for some time and sampled uniformly from this memory bank to update the network
    - Not quite clear how this works

## Double Q-learning

- Note that the max operator in DQN uses the same values both to select and to evaluate an action
  - We choose the  $a$  that maximizes the  $Q$  and take this same value of  $Q$  to construct the target towards which we tune  $Q$ .
  - What if we chose  $a$  according to some  $Q$  but then use another  $Q$  to evaluate  $Q(s,a)$  for the  $a$  that we found ?
  - Basically it means doing the following
    - In Q learning (still not difference from above, just a clearer notation):
$$Y_t^Q = R_{t+1} + \gamma Q(S_{t+1}, \arg\max_a Q(S_{t+1}, a; \theta_t); \theta_t)$$
    - And now apply the separation to get the double Q learning. Note the different set of parameters:
$$Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \arg\max_a Q(S_{t+1}, a; \theta_t); \theta_t')$$
  - This second set of weights can be updated symmetrically by switching the roles of  $\theta$  and  $\theta'$

## Double DQN

- We already have a separate network for the target network in DQN. Can we use this target network to evaluate the action and the online network to choose it ?
- i.e. instead of having another network  $\theta'$  for the double Q, use the target network  $\theta^-$  that we used in the DQN and get the double DQN :
$$Y_t^{\text{DoubleDQN}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \arg\max_a Q(S_{t+1}, a; \theta_t), \theta_t^-)$$
- The update of the target network remains the same as in DQN (periodically, once every few frames, copy weights from the online network)